

**Ph.D. Defense**

**Meta-Reinforcement Learning with Imaginary Tasks**

Suyoung Lee

Advisor: Youngchul Sung

Smart Information Systems Research Lab (SISReL)  
School of Electrical Engineering, KAIST

# Introduction

- Deep reinforcement learning (RL) faces critical challenges, including the issues of **overfitting** and a **limited generalization** capability when encountering **unseen test tasks**.
- Meta-reinforcement learning (meta-RL) aims to solve these problems by training agents across a variety of tasks, enabling them to **learn to infer** the underlying dynamics of new tasks and to **rapidly adapt** their policies accordingly.
  - Nonetheless, conventional meta-RL methods rely on a restricted training task distribution, which limits adaptability to **out-of-distribution** (OOD) test tasks.
- This thesis introduces two meta-RL algorithms that train the policy on **imaginary tasks** generated by a learned dynamics model, enhancing the ability to generalize to unseen distribution of tasks.

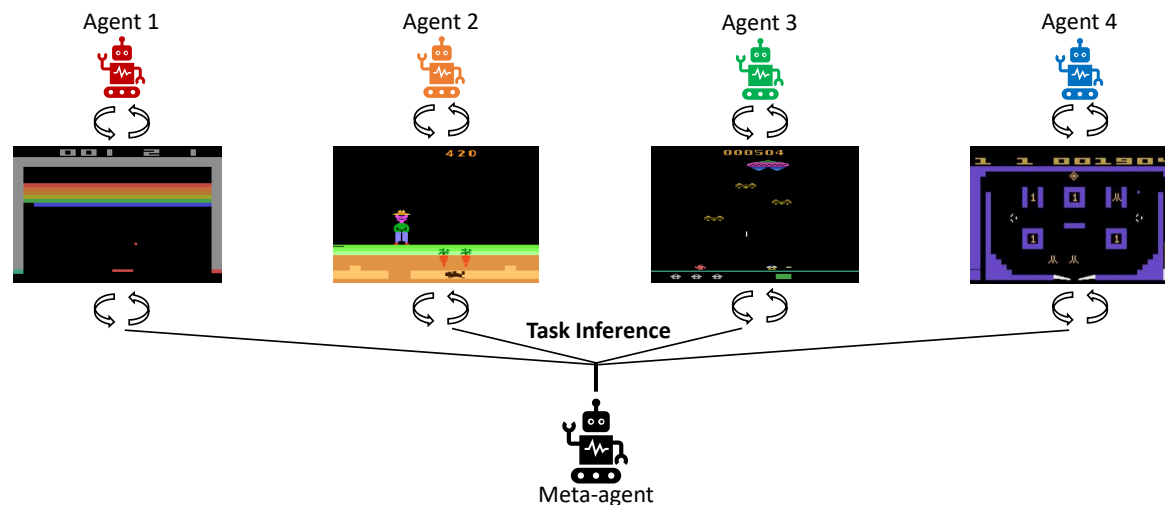


Figure 1: Comparison of standard RL and meta-RL.

# Outline and Contributions

---

## Meta-Reinforcement Learning with Imaginary Tasks

1. Introduction to Meta-Reinforcement Learning: A running example on Gridworld
2. Latent Dynamics Mixture (LDM)<sup>1</sup>
  - Addresses **parametric** task variability (Gridworld, MuJoCo).
  - Trains a policy on imaginary tasks generated from **interpolations** of latent beliefs.
  - Enhances generalization to unseen out-of-distribution (OOD) tasks with **simple parametric variations**.
3. Subtask Decomposition and Virtual Training (SDVT)<sup>2</sup>
  - Addresses **non-parametric** task variability (Meta-World).
  - Learns to **decompose** each non-parametric task into a set of shared **elementary subtasks**.
  - Trains a policy on imaginary tasks composed of **imaginary compositions** of learned elementary subtasks.

---

<sup>1</sup>Suyoung Lee and Sae-Young Chung, “Improving Generalization in Meta-RL with Imaginary Tasks from Latent Dynamics Mixture,” NeurIPS 2021.

<sup>2</sup>Suyoung Lee, Myungsik Cho, and Youngchul Sung, “Parametrizing Non-Parametric Meta-Reinforcement Learning Tasks via Subtask Decomposition,” NeurIPS 2023.

# Meta-Reinforcement Learning

---

- Assume a finite task space:  $\mathcal{M} = \{M^{(1)}, \dots, M^{(m)}\}$ .
- Meta-reinforcement learning (Meta-RL) aims to **learn to adapt** to a **set of** Markov Decision Processes (MDPs) with diverse reward and transition dynamics.

$$M^{(i)} = (\mathcal{S}, \mathcal{A}, R^{(i)}, T^{(i)}, T_0^{(i)}, \gamma, H). \quad (1)$$

- At the start of a **meta-episode**, an MDP  $M^{(i)} \sim \mathbb{P}(\mathcal{M})$  is sampled from the pool of tasks  $\mathcal{M}$ .
- Each MDP  $M^{(i)}$  has unique dynamics:
  - **Reward** dynamics:  $R^{(i)}(r_{t+1} \mid s_t, a_t, s_{t+1})$  or  $r_{t+1} = R^{(i)}(s_t, a_t, s_{t+1})$  (deterministic case)
  - **(State) Transition** dynamics:  $T^{(i)}(s_{t+1} \mid s_t, a_t)$
  - **Initial state** distribution:  $T_0^{(i)}(s_0)$
- A **meta-episode** consists of  $N$  **rollout episodes**, each with a horizon  $H$ .
  - A new task  $M^{(i)}$  is sampled only at the start of each **meta-episode** (every  $H^+ := NH$  steps).
  - After each **rollout episode**, the state is reset to an initial state (every  $H$  steps).

$$s_{kH} \sim T_0^{(i)}(\cdot), \quad k = 0, 1, \dots, N - 1. \quad (2)$$

## Meta-Reinforcement Learning

---

- A belief is defined as a posterior distribution over the reward and state transition functions given the current meta-episode's trajectory  $\tau_{:t} = (s_0, a_0, r_1, s_1, a_1, r_2, \dots, s_{t-1}, a_{t-1}, r_t, s_t)$ .

$$b_t \left( R^{(i)}, T^{(i)} \right) := \mathbb{P}_{R,T} \left( R^{(i)}, T^{(i)} \mid \tau_{:t} \right), \quad i = 1, \dots, m. \quad (3)$$

- The optimization of meta-RL typically consists of two stages.
  - **Task Inference**: Learn to infer the current task by constructing a belief.
  - **Policy Optimization**: Learn the optimal policy conditioned on the inferred belief.
- The objective is to optimize the belief encoder  $b_t = q_\phi(\tau_{:t})$  and the policy  $\pi_\psi(a_t \mid s_t, b_t)$  that maximize the expected return across all MDPs.

$$J(\phi, \psi) = \mathbb{E}_{M^{(i)} \sim \mathbb{P}(\mathcal{M})} \left[ \mathbb{E}_{T_0^{(i)}, T^{(i)}, \pi} \left[ \sum_{t=0}^{H^+-1} \gamma^t R^{(i)}(s_t, a_t, s_{t+1}) \right] \right]. \quad (4)$$

## Running Example

- Gridworld navigation task ( $3 \times 3$ ) without obstacles.




Figure 2:  $3 \times 3$  Gridworld navigation task  $M^{(7)}$  with an unknown goal position  $\times$ .

- $\mathcal{M} = \{M^{(1)}, \dots, M^{(8)}\}$ .
- $\mathcal{S} = \{x_0, \dots, x_8\}$  and  $\mathcal{A} = \{\text{'up,' 'down,' 'left,' 'right,' 'stay'}\}$ .
- All  $M^{(i)}$ 's share the same standard Gridworld transition  $T^{(i)} = T$ .
- All  $M^{(i)}$ 's start from the same initial state  $x_0$ . i.e.,  $T_0^{(i)}(x_0) = T_0(x_0) = 1$ .
- Each task  $M^{(i)}$  assigns a reward of 1 only when the agent reaches  $x_i$ .

$$R^{(i)}(s_t, a_t, s_{t+1}) = \begin{cases} 1 & \text{if } s_{t+1} = x_i \\ 0 & \text{otherwise} \end{cases} \quad (5)$$

- $N = 2$  rollout episodes,  $H = 10$  steps.

## Running Example

- **Meta-episode 1:** Suppose  $M^{(7)}$  is sampled for the first meta-episode.




Figure 3: Example trajectory in meta-episode 1 within  $M^{(7)}$ .

- **Meta-episode 2:** After  $H^+ = NH = 2 \times 10$  steps, a new task  $M^{(5)}$  is sampled.




Figure 4: Example trajectory in meta-episode 2 within  $M^{(5)}$ .





## Belief Update Example

---

- Belief update example of a perfectly trained meta-RL agent.




Figure 6: A meta-episode on task  $M^{(7)}$ . The first rollout episode until  $t = 2$ .

- ( $t = 1$ ) Assume we take action  $a_0 = \text{'right'}$  at  $s_0$ . We observe  $s_1 = x_1$  and since the current meta-episode's task is  $M^{(7)}$ , we observe  $r_1 = 0$ , implying  $\mathbb{P}_R(R^{(1)}|\tau_{:1}) = 0$ .

$$\left\{ b_1 \left( R^{(i)} \right) \right\}_{i=1}^8 = \left\{ 0, \frac{1}{7}, \frac{1}{7}, \frac{1}{7}, \frac{1}{7}, \frac{1}{7}, \frac{1}{7}, \frac{1}{7} \right\}. \quad (8)$$

- ( $t = 2$ ) Continuing with action  $a_1 = \text{'right'}$  at  $s_1$ , observing  $s_2 = x_2$  and  $r_2 = 0$  refines the belief further.

$$\left\{ b_2 \left( R^{(i)} \right) \right\}_{i=1}^8 = \left\{ 0, 0, \frac{1}{6}, \frac{1}{6}, \frac{1}{6}, \frac{1}{6}, \frac{1}{6}, \frac{1}{6} \right\}. \quad (9)$$

## Belief Update Example

- Belief update example of a **perfectly trained meta-RL agent**.




Figure 7: A meta-episode on task  $M^{(7)}$ . The first rollout episode until  $t = 5$ .

- ( $t = 4$ ) Taking action  $a_3 = \text{'up'}$  at  $s_3$  leads to  $s_4 = x_8$  and  $r_4 = 0$ , refining the belief.

$$\left\{ b_4 \left( R^{(i)} \right) \right\}_{i=1}^8 = \left\{ 0, 0, \frac{1}{4}, \frac{1}{4}, 0, \frac{1}{4}, \frac{1}{4}, 0 \right\}. \quad (10)$$

- ( $t = 5$ ) Finally, the action  $a_4 = \text{'left'}$  at  $s_4$  leads to the goal, observing  $s_5 = x_7$  and  $r_5 = 1$ .

$$\left\{ b_5 \left( R^{(i)} \right) \right\}_{i=1}^8 = \{ 0, 0, 0, 0, 0, 0, 1, 0 \}. \quad (11)$$

- We confirm that the current task of this meta-episode is certainly  $M^{(7)}$ .
- This conclusive belief remains for the rest of the meta-episode until a new task is sampled.

## Belief Update Example

- Belief update example of a **perfectly trained meta-RL agent**.




Figure 8: A meta-episode on task  $M^{(7)}$ . The second rollout episode.

- After the end of the first rollout episode ( $H = 10$  steps), the state is reset to the initial state  $s_H = x_0$ .
- However, since the agent is in the same meta-episode, it is still in  $M^{(7)}$ , keeping the belief.

$$\left\{ b_t \left( R^{(i)} \right) \right\}_{i=1}^8 = \{0, 0, 0, 0, 0, 0, 1, 0\}, \quad \text{for } 5 \leq t < H^+. \quad (12)$$

- The policy can exploit the belief inferred from the first rollout episode.
- After the meta-episode terminates (i.e.,  $H^+ = 20$  steps), a new task is sampled for the next-meta-episode. Also, the belief is reset to the prior  $b_0$ .

## Bayes-Adaptive Meta-RL

---

- To learn  $b_t(R^{(i)}, T^{(i)}) := \mathbb{P}_{R,T}(R^{(i)}, T^{(i)} \mid \tau_{:t})$ , we need to learn the prior and posterior update.
  - Prior:  $b_0(R^{(i)}, T^{(i)}) := \mathbb{P}_{R,T}(R^{(i)}, T^{(i)} \mid s_0)$ .
  - Posterior update upon taking  $a_t$  from  $s_t$  and observing  $s_{t+1}$  and  $r_{t+1}$  for all  $t = 0, \dots, H^+ - 1$ .

$$b_{t+1}(R^{(i)}, T^{(i)}) = \frac{\mathbb{P}(s_{t+1}, r_{t+1} \mid s_t, a_t, R^{(i)}, T^{(i)})}{\sum_{k=1}^m \mathbb{P}(s_{t+1}, r_{t+1} \mid s_t, a_t, R^{(k)}, T^{(k)}) b_t(R^{(k)}, T^{(k)})} b_t(R^{(i)}, T^{(i)}), \quad (13)$$

where the likelihood  $\mathbb{P}(s_{t+1}, r_{t+1} \mid s_t, a_t, R^{(i)}, T^{(i)}) = R^{(i)}(r_{t+1} \mid s_t, a_t, s_{t+1})T^{(i)}(s_{t+1} \mid s_t, a_t)$ .

Details in the Appendix (Eq. 27)

- The Bayes-adaptive policy  $\pi(a_t \mid s_t, b_t)$  is trained to maximize the return conditioned on the belief.
- The posterior update requires a tractable task space and knowledge of the reward and transition functions.

# Motivation




Figure 9: An example of test trajectory that is trained on  $\mathcal{M}_{\text{train}} = \{M^{(i)}\}_{i=1}^8 - M^{(4)}$  and tested on  $\mathcal{M}_{\text{test}} = M^{(4)}$ .

- Traditional meta-RL methods operate under the assumption that training and test tasks are drawn from the **same distribution**.

$$\mathcal{M} = \mathcal{M}_{\text{train}} = \mathcal{M}_{\text{test}}.$$

- This research is prompted by the following motivating question.

What if we evaluate the agent on *out-of-distribution (OOD)* test tasks?

$$\mathcal{M} = \mathcal{M}_{\text{train}} \cup \mathcal{M}_{\text{test}} \quad \text{and} \quad \mathcal{M}_{\text{train}} \cap \mathcal{M}_{\text{test}} = \emptyset.$$

Is it possible for the agent to effectively **explore and exploit** these previously unseen tasks?

# **Improving Generalization in Meta-RL with Imaginary Tasks from Latent Dynamics Mixture**

**Suyoung Lee** and Sae-Young Chung  
Presented at NeurIPS 2021

# Introduction

- Similar to the Gridworld example, most conventional meta-RL methods assume the **same distribution** of training and test tasks:  $(\mathcal{M} = \mathcal{M}_{\text{train}} = \mathcal{M}_{\text{test}})$ .
- This research delves into OOD scenarios where the training and test tasks are made **completely disjoint**:  $\mathcal{M} = \mathcal{M}_{\text{train}} \cup \mathcal{M}_{\text{test}}$  and  $\mathcal{M}_{\text{train}} \cap \mathcal{M}_{\text{test}} = \emptyset$ .

Table 1: Set of training and test parameters for OOD MuJoCo tasks. Here,  $k \in \{0, 1, 2, 3\}$ .

	Ant-direction $\theta$	Ant-goal $r$	$\theta$	Half-cheetah-velocity $v$
$\mathcal{M}_{\text{train}}$	$90^\circ \times k$	$[0.0, 1.0) \cup [2.5, 3.0)$	$[0^\circ, 360^\circ)$	$[0.0, 0.5) \cup [3.0, 3.5)$
$\mathcal{M}_{\text{test}}$	$90^\circ \times k + 45^\circ$	$[1.0, 2.5)$	$[0^\circ, 360^\circ)$	$[0.5, 3.0)$




Figure 10: Illustrative examples of OOD MuJoCo tasks.

## Latent Dynamics Mixture: Key Idea

- **Step 1.** **Encode** the dynamics of each MDP into a latent space representation,  $z$ .
- **Step 2.** **Mix** these latent embeddings from multiple training tasks to form a mixture embedding  $\tilde{z}$ .
- **Step 3.** **Decode** this combined latent vector to synthesize an *imaginary* MDP,  $\tilde{M}$ , which is then used for training the policy network.




Figure 11: Conceptual flow of the Latent Dynamics Mixture (LDM) algorithm.



# Variational Bayes-Adaptive Deep RL

- We need an encoder and decoder to **infer** (parameterize) the task dynamics and to learn the dynamics model to **generate** an imaginary MDP.
- VariBAD<sup>3</sup> addresses the challenge of posterior updates, which are often computationally infeasible, by leveraging variational inference within meta-RL. The framework is structured as follows:
  - A Variational AutoEncoder (VAE) with a recurrent network  $q_\phi(\tau_{:t})$  that encodes the trajectory  $\tau_{:t}$  into a latent representation.
  - The decoders  $p_{\theta_R}$  and  $p_{\theta_T}$  that reconstruct the reward and transition dynamics.
  - A separate policy network  $\pi_\psi(a_t|s_t, b_t)$  conditioned on the inferred belief from the encoder.

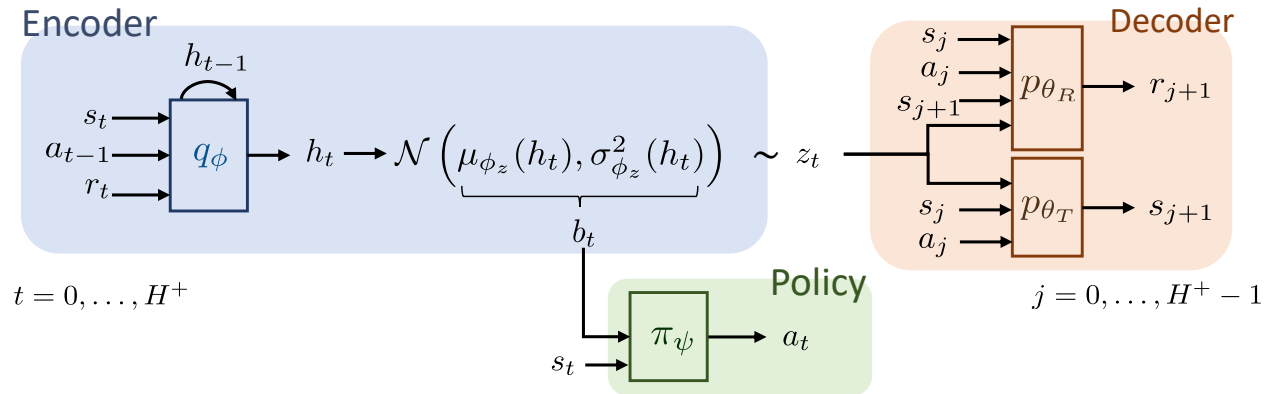


Figure 12: VariBAD architecture.

<sup>3</sup>Zintgraf, et al., “VariBAD: A Very Good Method for Bayes-Adaptive Deep RL via Meta-Learning,” ICLR 2020.

# Variational Bayes-Adaptive Deep RL

- The latent mean  $\mu_{\phi_z}(h_t)$  and variance  $\sigma_{\phi_z}^2(h_t)$  of the VAE are neural network outputs given the context of the current meta-episode,  $h_t = q_\phi(\tau_{:t})$ .
- The belief  $b_t = (\mu_{\phi_z}(h_t), \sigma_{\phi_z}^2(h_t))$  is expected to contain the **inferred task dynamics** until time  $t$ .
- As  $t$  grows and the agent explores the task space,
  - the uncertainty  $\sigma_{\phi_z}^2(h_t)$  about the current meta-episode's task diminishes.
  - the converged mean  $\mu_{\phi_z}(h_t)$  represents the **inferred parametrization** of the current task.




Figure 13: VariBAD's belief update example on the Ant-direction task. The latent beliefs are 5-dimensional.

# Latent Dynamics Mixture

- Multiple workers with shared networks.
  - All workers jointly train a shared policy network and a latent dynamics network.
  - Unlike VariBAD, we use two **separate encoders**  $q_{\phi_p}$  and  $q_{\phi_v}$ .
  - The policy part is optimized with PPO, and the VAE is trained to maximize ELBO.

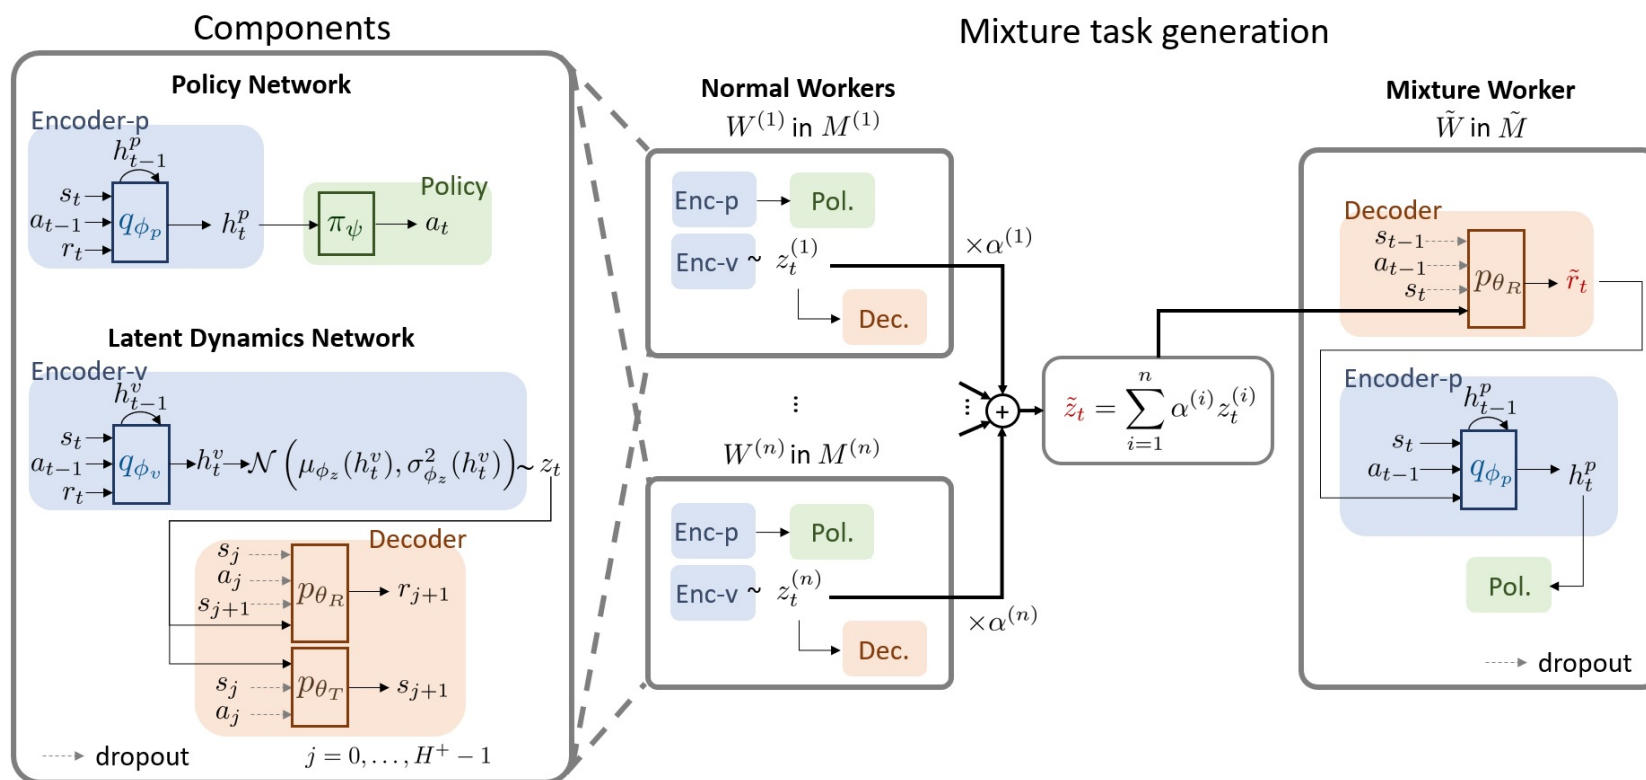
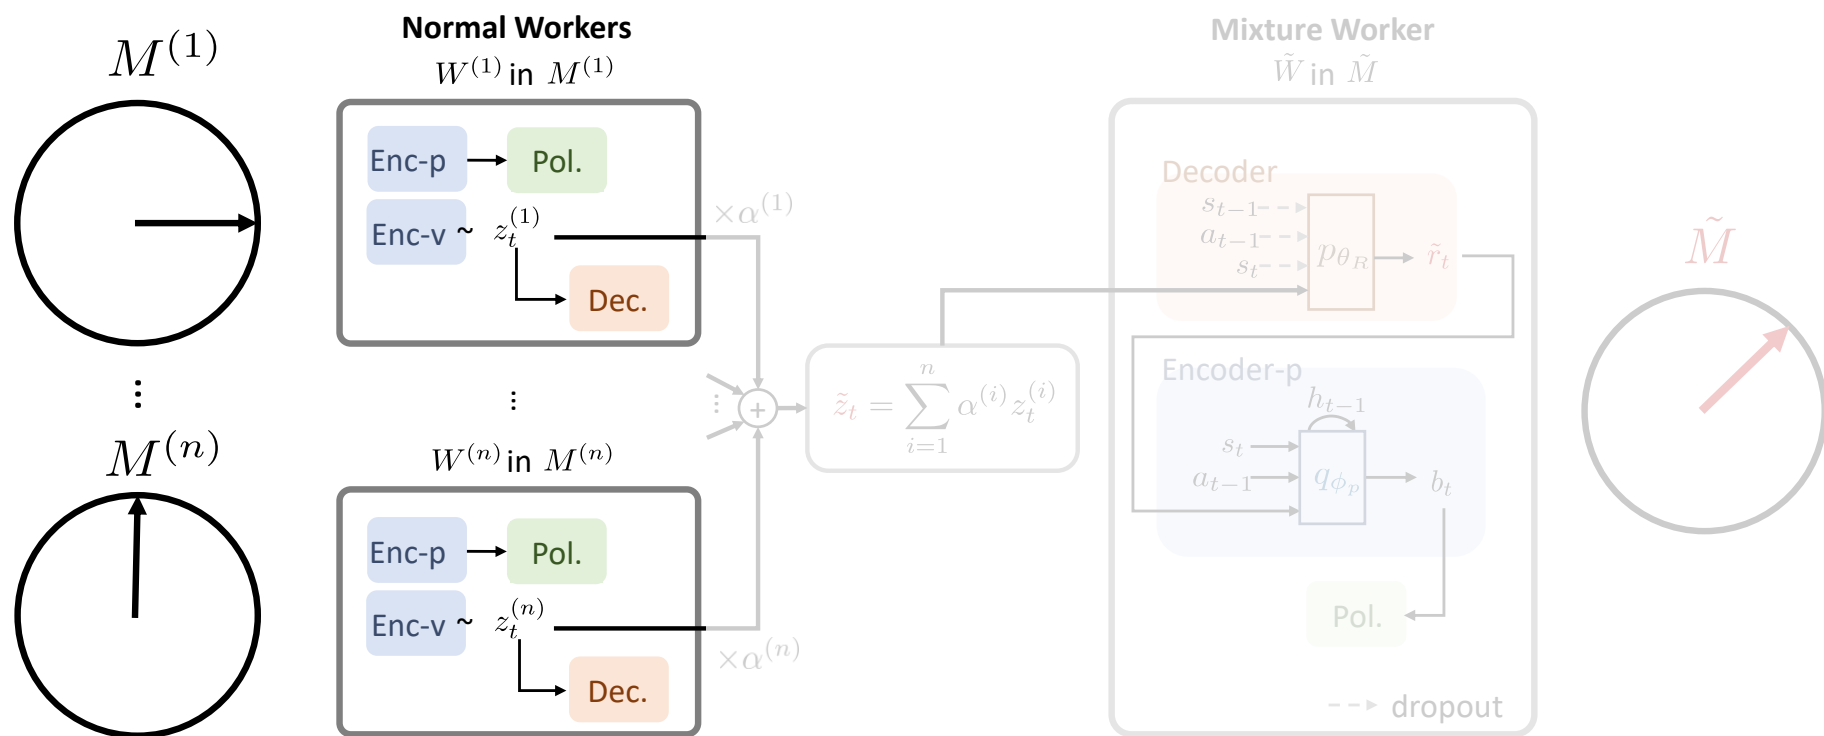


Figure 14: Imaginary task generation from Latent Dynamics Mixture (LDM).

# Latent Dynamics Mixture

- **Step 1.** Each **normal worker**  $W^{(i)}$  encodes its belief for  $M^{(i)}$  sampled from  $\mathbb{P}(\mathcal{M}_{\text{train}})$

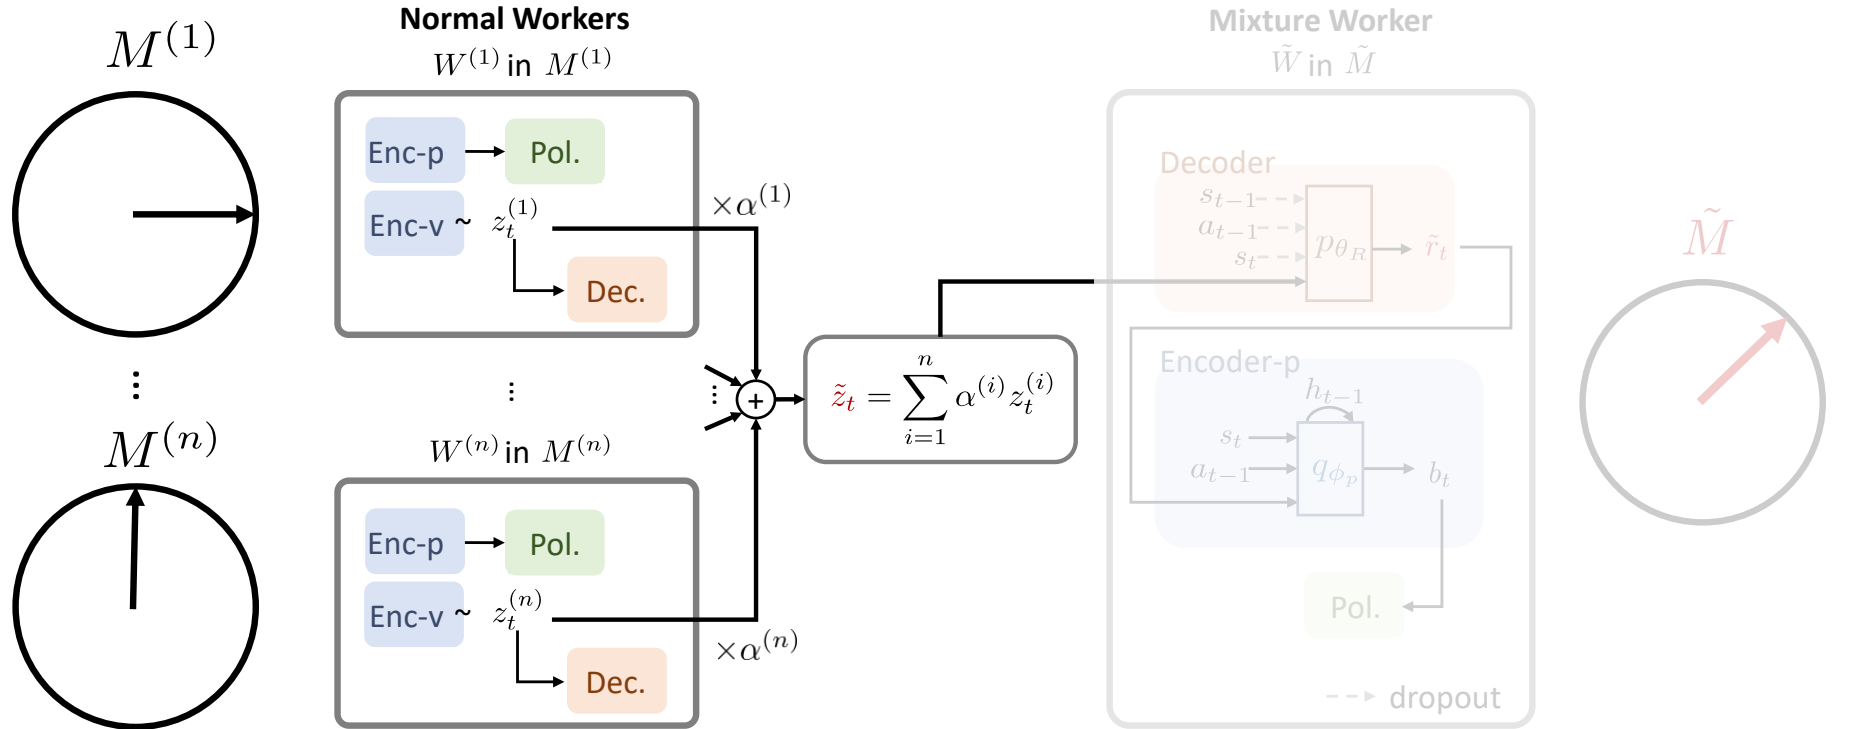


# Latent Dynamics Mixture

- **Step 2.** A **mixture worker** computes a mixture latent belief  $\tilde{z}_t = \sum_{i=1}^n \alpha^{(i)} z_t^{(i)}$ , where

$$\left(\alpha^{(1)}, \dots, \alpha^{(n)}\right) \sim \beta \cdot \text{Dirichlet}(\mathbb{1}_n) - \frac{\beta - 1}{n}. \quad (14)$$

- The weights satisfy  $\sum_{i=1}^n \alpha^{(i)} = 1$ ,  $\mathbb{E}[\alpha^{(i)}] = 1/n$ .
- The hyperparameter  $\beta$  controls the extrapolation level. If  $\beta = 1$ ,  $0 \leq \alpha^{(i)} \leq 1$ .

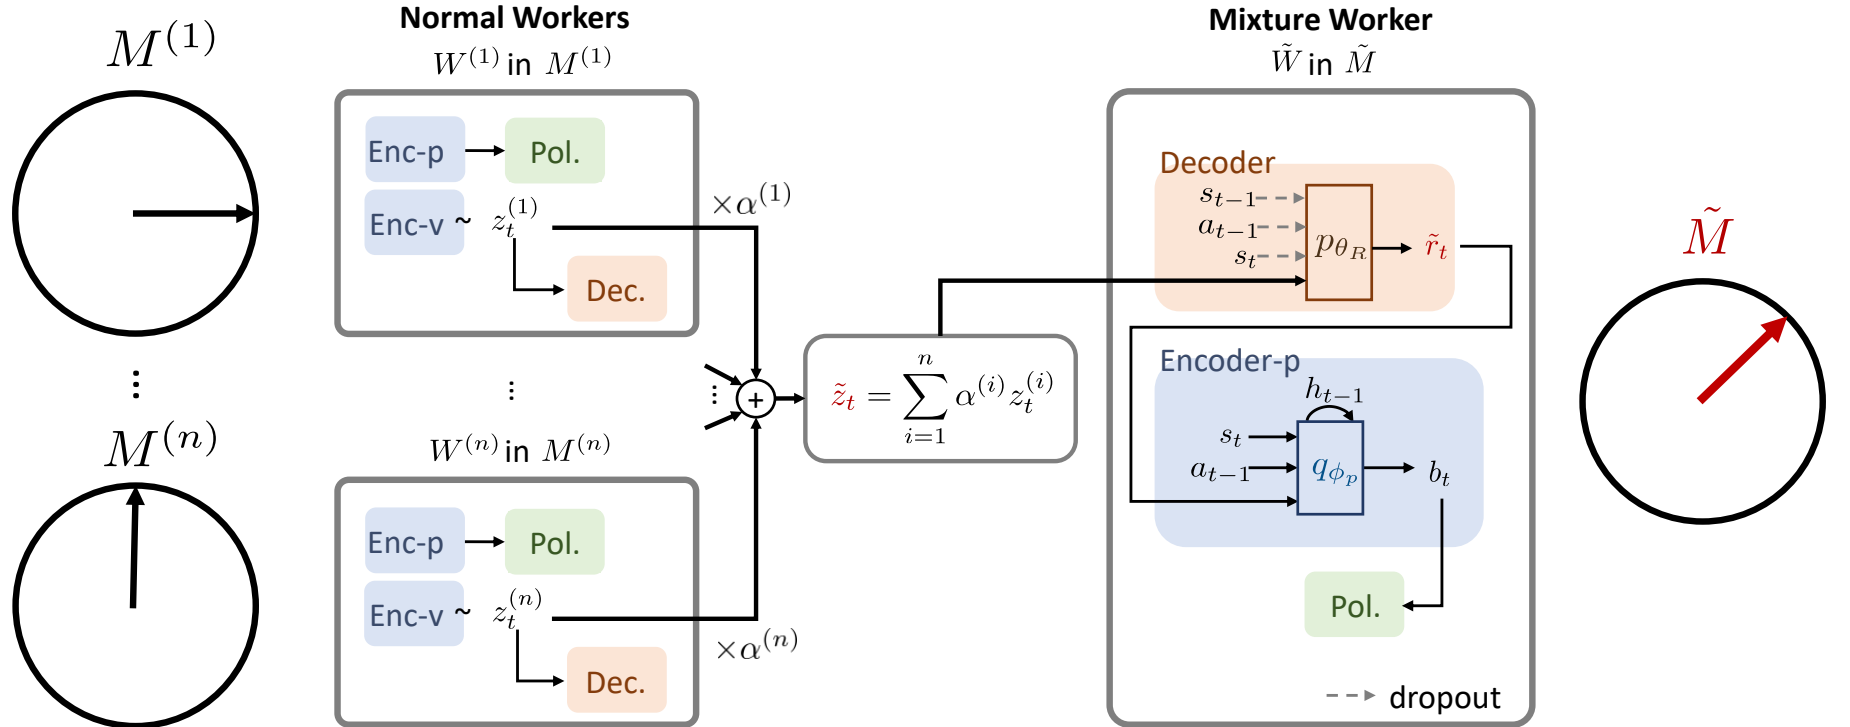


# Latent Dynamics Mixture

- **Step 3.** Given  $\tilde{z}_t$ , the mixture worker  $\tilde{W}$  generates an imaginary MDP  $\tilde{M}$  with imaginary rewards  $\tilde{r}_t \sim p_{\theta_R}(\cdot | s_t, a_t, s_{t+1}; \tilde{z}_t)$ . But the states are from a real training MDP  $M^{(k)} \sim \mathbb{P}(\mathcal{M}_{\text{train}})$ .

$$\tilde{M} = (\mathcal{S}, \mathcal{A}, \tilde{R}, T^{(k)}, T_0^{(k)}, \gamma, H). \quad (15)$$

- The trajectories from  $\tilde{M}$  are only used to train the policy part:  $q_{\phi_p}$  and  $\pi_{\psi}$ , not the VAE part.



## Dropout on Decoder Input

- The decoder easily overfits the state and action observations, **ignoring the latent belief  $z_t$** .
- States that are unseen during training may be assigned low rewards regardless of the mixture latent belief  $\tilde{z}_t$  (e.g., the dotted states in Figure 15).
- Therefore we apply **dropout** to all inputs of the decoder except the latent belief when training and generating with the decoder.

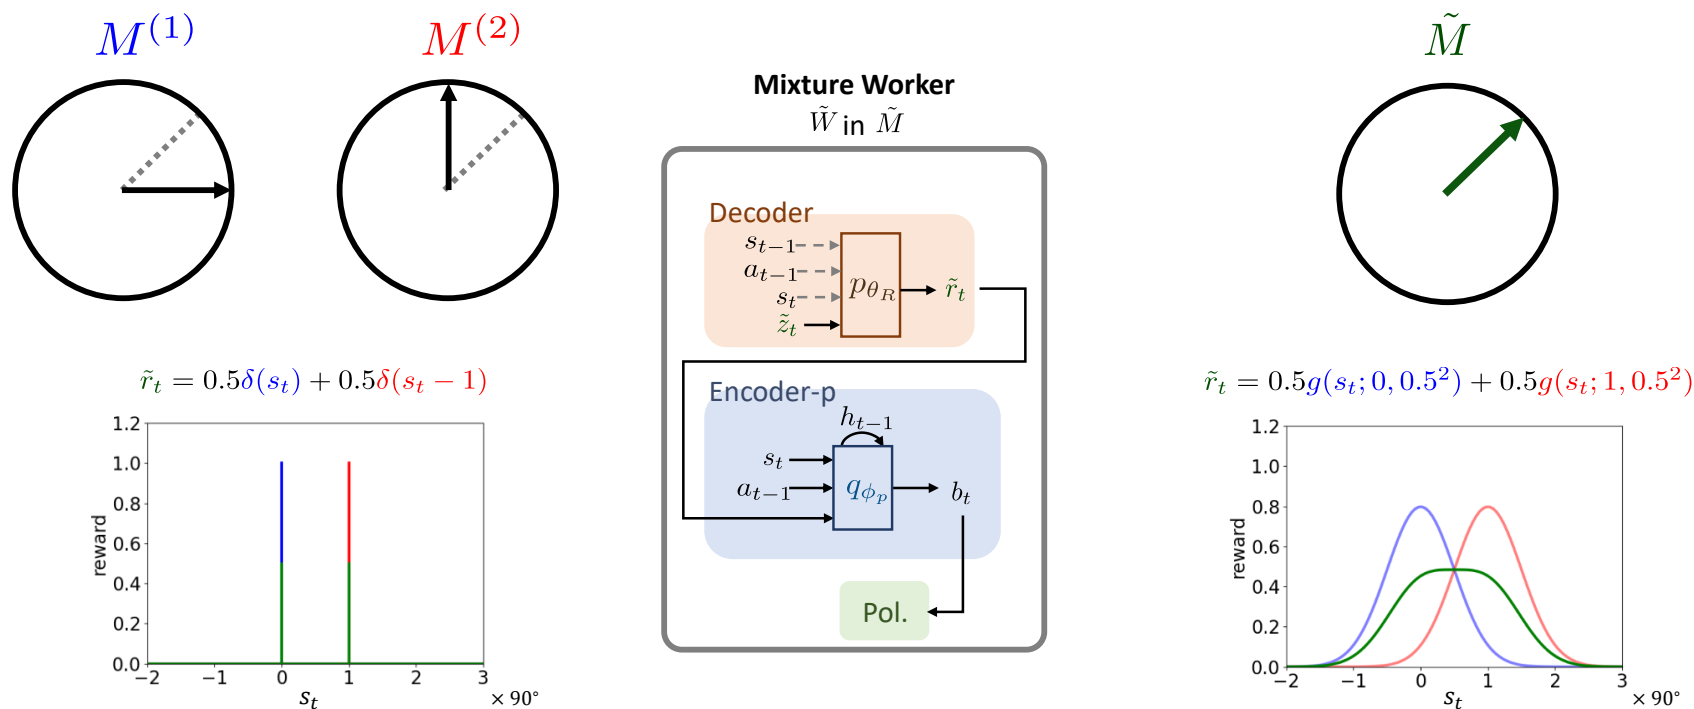


Figure 15: Applying dropout to the decoder input.

## Experiments – Setup

- We evaluate LDM and baselines on Gridworld and MuJoCo tasks, where we strictly divide the entire task distribution into disjoint training and test tasks.
- **Gridworld**
  - $7 \times 7$  Gridworld task with  $N = 4$  rollout episodes of horizon  $H = 30$  each.
  - Number of tasks:  $|\mathcal{M}_{\text{train}}| = 18$  and  $|\mathcal{M}_{\text{test}}| = 27$ .

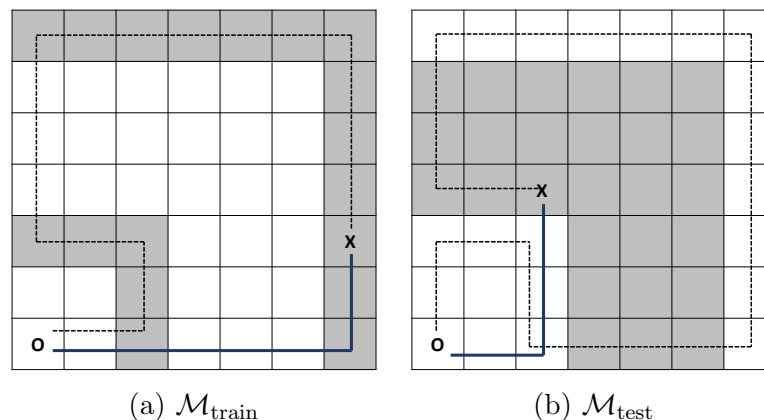


Figure 16: **Gridworld example.** (a) Training MDPs  $\mathcal{M}_{\text{train}}$ . (b) Test MDPs  $\mathcal{M}_{\text{test}}$ . A goal is located at one of the shaded positions.

- **MuJoCo**
  - We evaluate over three OOD MuJoCo tasks introduced in Figure 10 with  $N = 2$  and  $H = 200$ .



## Experiments – Gridworld

- We report the results at the last rollout episode in terms of
  - the mean returns in  $\mathcal{M}_{\text{train}}$  and  $\mathcal{M}_{\text{test}}$ .
  - the number of tasks in  $\mathcal{M}_{\text{test}}$  in which the agent fails to reach the goal (out of 27 tasks).
- The oracle methods, that are trained on the entire task set  $\mathcal{M} = \mathcal{M}_{\text{train}} \cup \mathcal{M}_{\text{test}}$  including the test tasks, are evaluated for reference.




Figure 17: Gridworld results.

## Experiments – Gridworld

- We provide an empirical analysis to confirm that LDM successfully generates **appropriate imaginary tasks** that contribute to solving the test tasks.
- The **dropout** applied to the next state input of the decoder plays a critical role here.
  - Dropout ( $p_{\text{drop}} = 0.7$ ) leads the decoder to attribute high rewards to certain goals in  $\mathcal{M}_{\text{test}}$ .

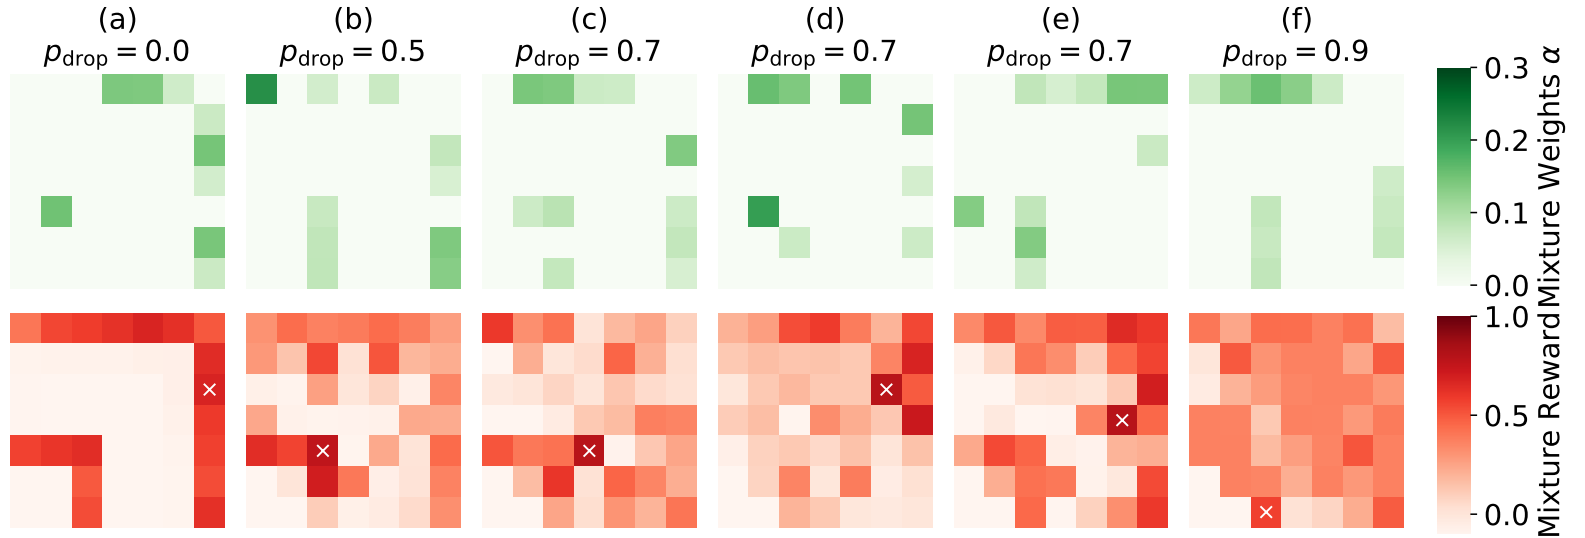


Figure 18: Examples of mixture tasks generated by LDM. **First row**: Mixture weights ( $\alpha^{(i)}$ ) multiplied to the latent beliefs ( $z_{H+}^{(i)}$ ). **Second row** (reward map): decoder output for each next state conditioned on the mixture weights from the first row.  $\times$  denotes the state yielding the maximum reward.

## Experiments – Gridworld

- We demonstrate that LDM’s applicability is not confined to target tasks within the interpolation of training tasks with the following experiment on Gridworld-extrapolation task.
  - Recall:  $(\alpha^{(1)}, \dots, \alpha^{(n)}) \sim \beta \cdot \text{Dirichlet}(\mathbb{1}_n) - \frac{\beta-1}{n}$ .
  - For lower  $\beta$  settings such as  $\beta = 1.0$  and  $\beta = 1.5$ , LDM concentrates on  $\mathcal{M}_{\text{test1}}$ .
  - As  $\beta$  increases, we observe a decline in returns for tasks in  $\mathcal{M}_{\text{test1}}$ , accompanied by an increase for tasks in  $\mathcal{M}_{\text{test2}}$ .

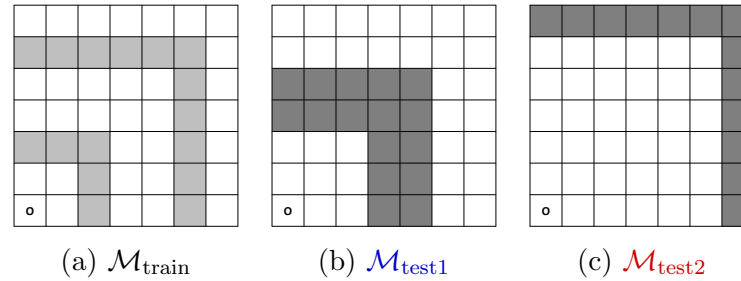


Figure 19: Gridworld-extrapolation task, that consists of two separate task regions,  $\mathcal{M}_{\text{test1}}$  and  $\mathcal{M}_{\text{test2}}$ .




Figure 20: Extrapolation results. Returns of LDM for different extrapolation level  $\beta$  on the Gridworld-extrapolation task.

## Experiments – MuJoCo

- We evaluate our method and baselines on three OOD meta-RL benchmarks: Ant-direction, Ant-goal, and Half-cheetah-velocity.
- The mean test returns in  $\mathcal{M}_{\text{eval}}$ , which is a fixed subset of  $\mathcal{M}_{\text{test}}$ , are illustrated in Figure 21.
- LDM surpasses the performance of non-oracle baselines.




Figure 21: MuJoCo results. Mean returns at the last rollout episode in  $\mathcal{M}_{\text{eval}}$ .

## Experiments – MuJoCo

- We demonstrate the sample trajectories of the agents in  $\mathcal{M}_{\text{eval}}$  at the last rollout episode.

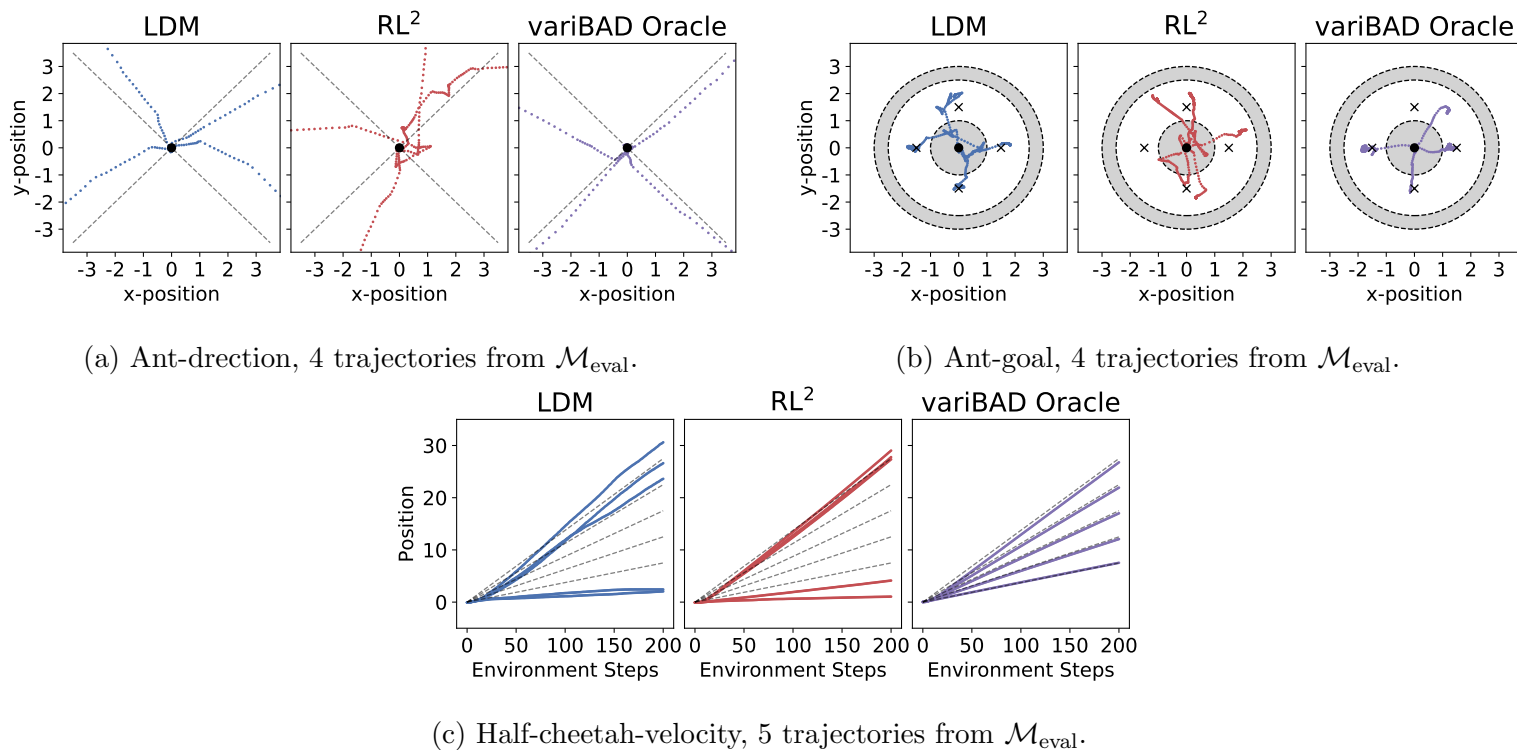


Figure 22: Sampled trajectories of the agents in  $\mathcal{M}_{\text{eval}}$ . The targets of  $\mathcal{M}_{\text{eval}}$  are indicated as dashed lines or cross marks

## Analysis on Task Embedding

- We have conducted empirical tests that affirm the latent models effectively reflect the **structure of the test task**, although not trained in  $\mathcal{M}_{\text{test}}$ , which supports the efficacy of LDM.
  - We sample 48 tasks in Ant-goal as in Figure 23a. 32 tasks from  $\mathcal{M}_{\text{train}}$  and 16 from  $\mathcal{M}_{\text{test}}$ .
  - For each task, We evaluate latent beliefs  $z_{H+}$  at the end of each meta-episode.

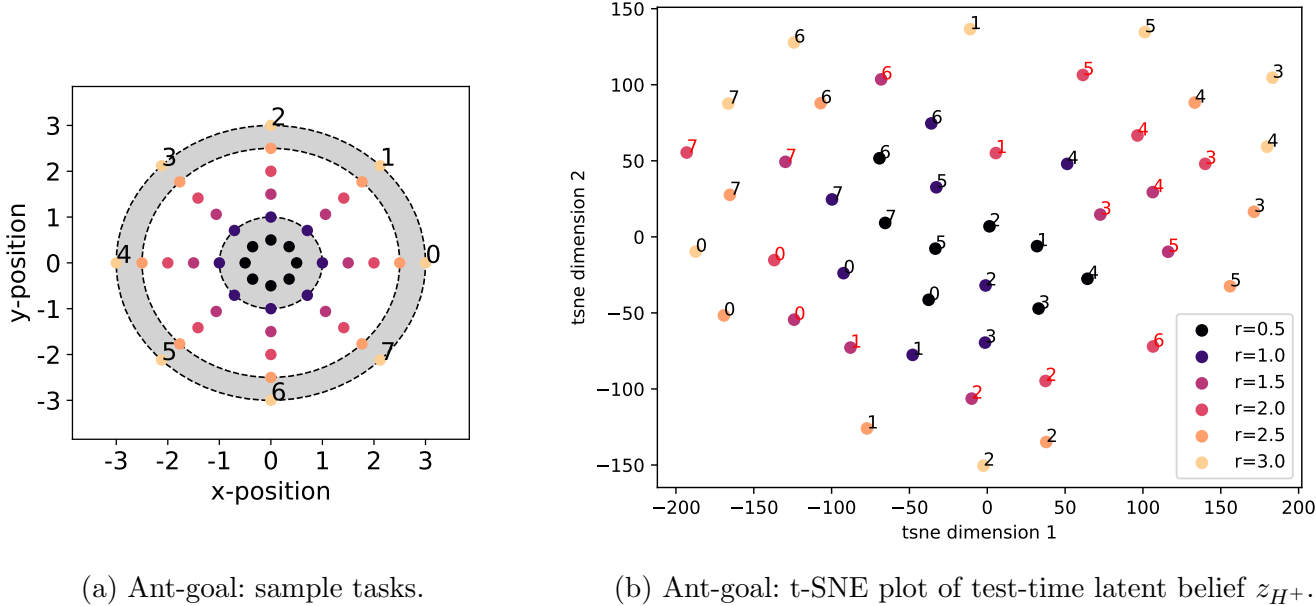


Figure 23: Latent belief distributions. Latent dynamics network’s learned latent models on Gridworld and Ant-goal. The red numbers denote the tasks that belong to  $\mathcal{M}_{\text{test}}$ .

## Limitation

---

- Parametric task variability
  - LDM requires the tasks to exhibit **simple parametric variations**, such as the goal position, target direction, and target velocity, that could be easily modeled by the latent embedding.
  - If the task distribution is more complex, we can not expect the test tasks to be represented as the **interpolations** of the latent embedding of the training tasks.
- Interpolation
  - LDM can generalize to test tasks that are within the **convex hull** of the training tasks' latent embeddings.
  - The extrapolation achieved by the hyperparameter  $\beta$  is limited.
- Tasks with varying reward dynamics
  - LDM prepares for test tasks that exhibit **unseen reward dynamics**.
  - It uses states from the real training tasks, therefore unable to prepare for OOD tasks with varying **state transition** dynamics.

# Parameterizing Non-Parametric Meta-Reinforcement Learning Tasks via Subtask Decomposition

**Suyoung Lee**, Myungsik Cho, and Youngchul Sung  
To be presented at NeurIPS 2023



## Introduction

- Conventional Meta-RL methods and LDM are evaluated on tasks with variations that can be expressed in a shared **parametric** form representing the task dynamics (Figure 24a and 24b).
- Now, we explore a more general meta-RL framework that addresses generalization in more **qualitatively distinct** tasks, namely with **non-parametric task variability**<sup>4</sup>.
  - For example, training on the task “Pick-place” and testing on “Sweep-into” from Meta-World (Figure 24c).




Figure 24: Illustrating different meta-RL scenarios.  $\circ$ : object,  $\times$ : goal.

<sup>4</sup>Yu, et al., “Meta-World: A Benchmark and Evaluation for Multi-Task and Meta Reinforcement Learning,” CoRL 2019.

## Motivation

- Despite the complex non-parametric task variability, tasks may share common elementary subtasks.
  - “**Pick-place**” = “grip-object” + “place-object”
  - “**Sweep-into**” = “grip-object” + “push-object”
- Our primary strategy is to

*Decompose each non-parametric task into a set of shared elementary subtasks,*

which allows to parameterize a task with embeddings representing the **subtask composition** along with the conventional **parametric variations**.




Figure 25: Meta-World ML-10 robotic manipulation benchmark. There are 10 training tasks and 5 test tasks with non-parametric task variability. Within each task, there are 50 parametric variations.

## Subtask Decomposition and Virtual Training: Key Idea

---

- For example, we want
  - Embedding(“**Pick-place**”) = (0.5, 0.5, 0.0)
  - Embedding(“**Sweep-into**”) = (0.0, 0.5, 0.5)
  - Where each dimension of the embedding corresponds to the weights for the following subtasks: “place-object,” “grip-object,” and “push-object”.
- The key problem is to learn the
  - Basis: the set of elementary subtasks
  - Coefficients: the decomposition of each task given the set of elementary subtasks.
- We employ meta-learning for the subtask decomposition (**SD**) process using a Gaussian mixture variational autoencoder (GMVAE) that encodes the trajectory up to the current timestep into
  - Compositional information: latent categorical contexts
  - Parametric variability information: latent Gaussian contexts.
- To further enhance generalization to unseen compositions of learned elementary subtasks, we propose a virtual training (**VT**) process similar to that of LDM.

## Subtask Decomposition

- The proposed architecture of SDVT incorporates three main components: **encoder**, **decoder**, and **policy**, similar to the VariBAD architecture.
- The major difference is that the encoder involves a **categorical latent variable**  $y_t$ .

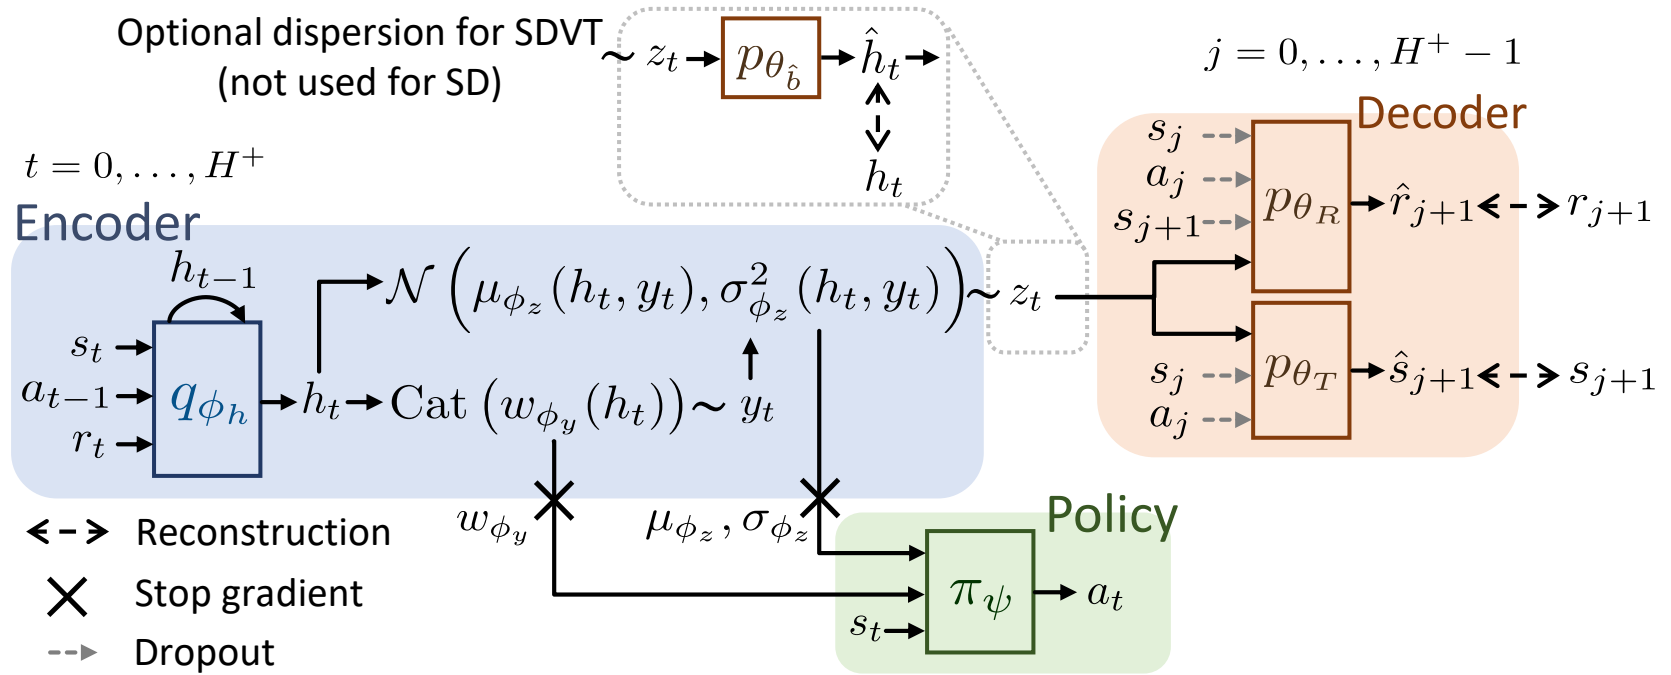


Figure 26: SDVT architecture.

## Subtask Decomposition

- Encoder:  $q_\phi(y_t, z_t|h_t) = q_{\phi_y}(y_t|h_t)q_{\phi_z}(z_t|h_t, y_t)$ 
  - A recurrent network encodes the past trajectory  $\tau_{:t}$  into a hidden embedding  $h_t = q_{\phi_h}(\tau_{:t})$ .
  - The categorical encoder  $q_{\phi_y}(y_t|h_t) : \text{Cat}(\omega_{\phi_y}(h_t))$  samples  $y_t$ , where  $\omega_{\phi_y}(h_t) \in \Delta^K$ .
  - Then the multivariate Gaussian encoder  $q_{\phi_z}(z_t|h_t, y_t) : \mathcal{N}(\mu_{\phi_z}(h_t, y_t), \sigma_{\phi_z}^2(h_t, y_t))$  samples a continuous latent context  $z_t$ .
  - $z_t$  contains the **compositional** information of the current task and the **parametric** information of the subtasks.

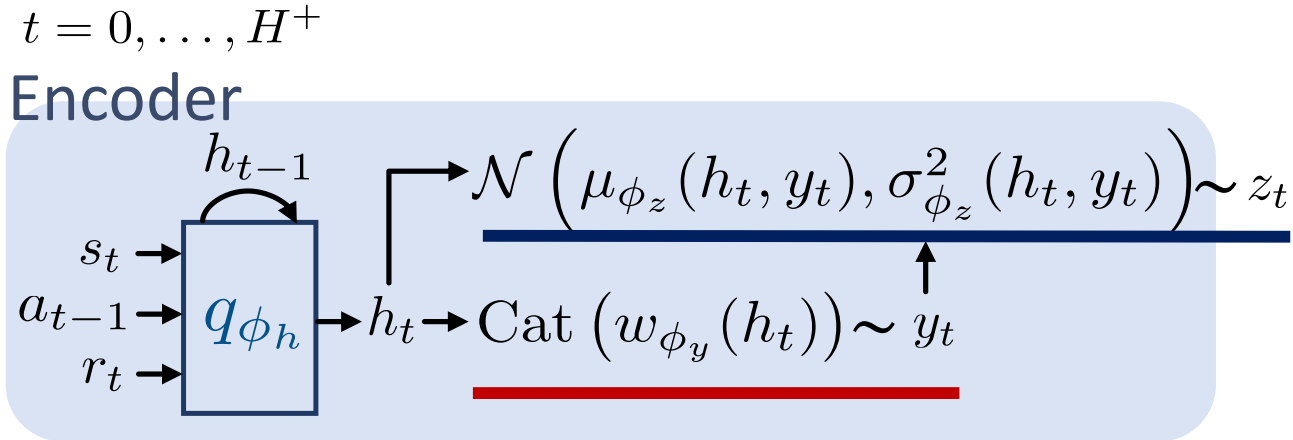


Figure 27: SDVT encoder.

## Learned Subtask Compositions

- Learned subtask compositions on ML-10.
  - Each column denotes the tasks in ML-10 ( $\mathcal{M}_{\text{train}} : 1 \sim 10, \mathcal{M}_{\text{test}} : 11 \sim 15$ ).
  - Each row number represents the index of the elementary subtask.
- We find that such learned subtask compositions  $y_{H^+} \in \Delta^K$  are shared by qualitatively similar tasks.
  - (3) “Pick-place,” (7) “Peg-insert-side,” (10) “Basketball”
  - (1) “Reach,” (5) “Drawer-close,” (8) “Window-open”




Figure 28: Learned subtask compositions on ML-10 by SDVT-LW with  $K = 5$ ,  $\alpha_c = 0.5$ .

## GMVAE Objectives

---

- The objective of the GMVAE is to maximize the evidence lower bound (ELBO), for  $t = 0, \dots, H^+$  and for the trajectory distribution  $d(M^{(k)}, \tau_{:H^+})$  at MDP  $M^{(k)}$ , induced by the policy  $\pi_\psi$ .

$$\text{ELBO}_t(\phi, \theta) = \mathbb{E}_{d(M^{(k)}, \tau_{:H^+})} \left[ \mathbb{E}_{q_\phi(y_t, z_t | h_t)} \mathcal{J}_{\text{GMVAE}} \right], \quad (16)$$

$$\mathcal{J}_{\text{GMVAE}} = \alpha_R \mathcal{J}_{\text{R-rec}} + \alpha_T \mathcal{J}_{\text{T-rec}} + \alpha_g \mathcal{J}_{\text{reg}} + \alpha_c \mathcal{J}_{\text{cat}}. \quad (17)$$

- Reconstruction objectives

$$\mathcal{J}_{\text{R-rec}} = \sum_{j=0}^{H^+-1} \log p_{\theta_R}(r_{j+1} | s_j, a_j, s_{j+1}; z_t), \quad (18)$$

$$\mathcal{J}_{\text{T-rec}} = \sum_{j=0}^{H^+-1} \log p_{\theta_T}(s_{j+1} | s_j, a_j; z_t). \quad (19)$$

- The reconstruction objectives are computed for all timesteps in the meta-episode, including the future ( $j > t$ ) and past ( $j < t$ ) (same as VariBAD and LDM).

## GMVAE Objectives

---

- The objective of the GMVAE is to maximize the evidence lower bound (ELBO), for  $t = 0, \dots, H^+$  and for the trajectory distribution  $d(M^{(k)}, \tau_{:H^+})$  at MDP  $M^{(k)}$ , induced by the policy  $\pi_\psi$ .

$$\text{ELBO}_t(\phi, \theta) = \mathbb{E}_{d(M^{(k)}, \tau_{:H^+})} \left[ \mathbb{E}_{q_\phi(y_t, z_t | h_t)} \mathcal{J}_{\text{GMVAE}} \right], \quad (20)$$

$$\mathcal{J}_{\text{GMVAE}} = \alpha_R \mathcal{J}_{\text{R-rec}} + \alpha_T \mathcal{J}_{\text{T-rec}} + \alpha_g \mathcal{J}_{\text{reg}} + \alpha_c \mathcal{J}_{\text{cat}}. \quad (21)$$

- Regularization objective

$$\mathcal{J}_{\text{reg}} = \log \frac{p_{\theta_z}(z_t | y_t)}{q_{\phi_z}(z_t | h_t, y_t)}. \quad (22)$$

- Unlike the standard VAE that assumes a **standard normal prior**, we learn  **$K$  distinct Gaussian priors** conditioned on  $y_t$  and minimize the KL divergence to the learned posterior  $q_{\phi_z}(z_t | h_t, y_t)$ .

- Categorical objective

$$\mathcal{J}_{\text{cat}} = \log \frac{p(y_t)}{q_{\phi_y}(y_t | h_t)}. \quad (23)$$


- The categorical objective  $\mathcal{J}_{\text{cat}}$  maximizes the **conditional entropy** of  $y_t$  given  $h_t$ .
- We use a fixed uniform prior  $p(y_t)$ .
- We set the coefficient  $\alpha_c$  large enough to penalize one-hot subtask compositions.



# Occupancy Regularization

- The **number of underlying subtasks**  $K$  (i.e., dimension of  $y_t$ ) is a crucial hyperparameter that should be determined based on the number of training tasks  $N_{\text{train}}$  and their similarities.
- We assume  $K^* < N_{\text{train}}$ , otherwise each task will be classified into a separate subtask with one-hot label, preventing learning shareable subtasks.
- We start with a sufficiently large  $K = N_{\text{train}}$  and regularize ELBO objective with the following **occupancy regularization** to progressively reduce the number of effective subtasks.

$$\mathcal{J}_{\text{occ}} = -\log K \left( e^{-K+1}, e^{-K+2}, \dots, e^{-1}, e^0 \right) \cdot y_t. \quad (24)$$

(a) SDVT ( $\alpha_o = 5.0$ ).(b) SDVT ( $\alpha_o = 10.0$ ).Figure 29: Occupancy ablation. Learned subtask compositions of SDVT ( $K = 10, \alpha_c = 1.0$ ) for different occupancy coefficients.

## Virtual Training

- LDM's task generation is limited to **parametric variations**.
- Now, using our GMVAE, we can condition the decoder on an **imaginary composition** of subtasks  $\tilde{y}$ .

$$\tilde{y} \sim \text{Dirichlet}(\bar{y}), \quad (25)$$

where  $\bar{y}$  is the empirical running mean of  $y_t$  aggregated during training.

- $\tilde{y}$  is fixed for a meta-episode, while  $\tilde{z}_t$  varies over time.
- We train the policy to maximize the sum of generated rewards  $\tilde{r}_{t+1} \sim p_{\theta_R}(\cdot | s_t, a_t, s_{t+1}; \tilde{z}_t)$ .






Figure 30: Generation of imaginary rewards with the decoder conditioned on a fixed imaginary subtask composition  $\tilde{y}$ .

## Virtual Training


- Generated imaginary tasks result in diverse trajectories.
- We evaluate the same policy conditioned on different subtask compositions  $\tilde{y}$ .
- All states are from the Meta-World “Reach.” Red rods: gripper and blue circles: object.




(a)  $\tilde{y} = [1, 0, 0, 0, 0]$ .



(b)  $\tilde{y} = [0, 1, 0, 0, 0]$ .



(c)  $\tilde{y} = [\frac{1}{2}, \frac{1}{2}, 0, 0, 0]$ .



(d)  $\tilde{y} = [0, 0, \frac{1}{3}, \frac{1}{3}, \frac{1}{3}]$ .

Figure 31: Trajectories on generated tasks.

## A Brief Summary

- SDVT = Subtask Decomposition (SD) and Virtual Training (VT).
- SDVT without a GMVAE and with a single Gaussian VAE reduces to LDM.
- LDM without virtual training reduces to VariBAD.
- VariBAD without a dynamics decoder reduces to RL<sup>2</sup>.

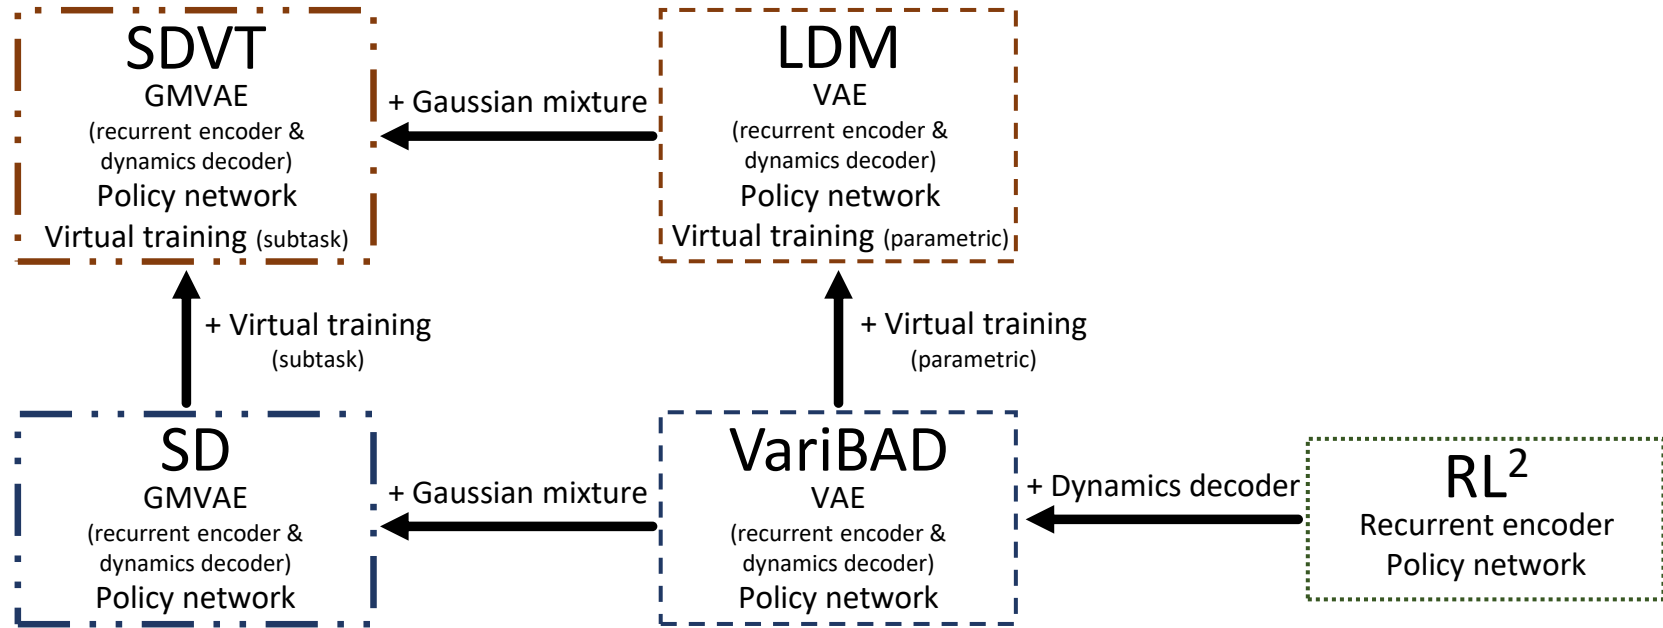


Figure 32: Schematic overview of proposed algorithms and baselines.

## Experiments – Setup

---

- The Meta-World V2 benchmark
  - A benchmark with 50 qualitatively distinct robotic manipulation tasks.
  - **ML-10** benchmark:  $N_{\text{train}} = 10$  training tasks and  $N_{\text{test}} = 5$  test tasks.
  - **ML-45** benchmark:  $N_{\text{train}} = 45$  training tasks and  $N_{\text{test}} = 5$  test tasks.
- Our methods
  - SDVT: subtask decomposition and virtual training.
  - SD: subtask decomposition **without virtual training**.
  - SDVT-LW and SD-LW (Light Weight variants): ours **without the occupancy regularization** ( $\alpha_o = 0$ ), assuming that we know the optimal number of elementary subtasks  $K^* = 5$ .  
The default setup is  $K = N_{\text{train}}$  with occupancy regularization ( $\alpha_o = 1$ ).
- Evaluation metric
  - $N = 10$  rollout episodes,  $H = 500$  steps per rollout episode.
  - We evaluate the performance at the **last rollout episode** in terms of the **success rate** and the **episode return**.

## Experiments – Result

- On the test tasks, SDVT and SDVT-LW substantially outperform all baselines.
- SDVT’s gain over LDM is attributed to our virtual training process, which is specifically designed for test tasks involving non-parametric variability.

Table 2: Meta-World V2 success rates and returns. We report the final success rates (%) and returns of our methods and baselines averaged across training tasks and test tasks of the ML-10 and ML-45 benchmarks. All results are reported as the mean success rate  $\pm$  95% confidence interval of 8 seeds.

Methods	Success Rate				Return			
	ML-10		ML-45		ML-10		ML-45	
	Train	Test	Train	Test	Train	Test	Train	Test
SDVT	<b>77.2<math>\pm</math>3.0</b>	32.8 $\pm$ 3.9	55.6 $\pm$ 4.2	28.1 $\pm$ 3.2	<b>3656<math>\pm</math>62</b>	1225 $\pm$ 160	2379 $\pm$ 214	839 $\pm$ 74
SDVT-LW	62.1 $\pm$ 4.1	<b>33.4<math>\pm</math>5.0</b>	50.4 $\pm$ 4.1	<b>31.2<math>\pm</math>1.2</b>	3454 $\pm$ 137	<b>1527<math>\pm</math>214</b>	2294 $\pm$ 202	<b>894<math>\pm</math>27</b>
SD	77.0 $\pm$ 5.9	30.8 $\pm$ 7.7	<b>61.0<math>\pm</math>1.7</b>	23.0 $\pm$ 5.1	3630 $\pm$ 241	1112 $\pm$ 190	<b>2672<math>\pm</math>79</b>	786 $\pm$ 69
SD-LW	75.5 $\pm$ 5.5	26.2 $\pm$ 8.7	56.7 $\pm$ 1.5	25.4 $\pm$ 2.9	3525 $\pm$ 297	1043 $\pm$ 234	2578 $\pm$ 64	793 $\pm$ 49
RL <sup>2</sup>	67.4 $\pm$ 4.4	15.1 $\pm$ 2.7	58.0 $\pm$ 0.4	11.8 $\pm$ 3.2	1159 $\pm$ 83	715 $\pm$ 33	1411 $\pm$ 22	663 $\pm$ 100
MAML	42.2 $\pm$ 4.5	3.9 $\pm$ 3.7	32.0 $\pm$ 1.4	19.8 $\pm$ 6.3	1822 $\pm$ 136	439 $\pm$ 78	1388 $\pm$ 104	658 $\pm$ 96
PEARL	23.2 $\pm$ 1.9	0.8 $\pm$ 0.5	10.3 $\pm$ 2.4	6.7 $\pm$ 3.3	1081 $\pm$ 77	340 $\pm$ 54	597 $\pm$ 121	506 $\pm$ 122
VariBAD	58.2 $\pm$ 8.9	14.1 $\pm$ 6.1	57.0 $\pm$ 1.2	22.1 $\pm$ 3.5	3055 $\pm$ 466	919 $\pm$ 143	2492 $\pm$ 47	762 $\pm$ 40
LDM	56.7 $\pm$ 12.3	19.8 $\pm$ 6.0	54.1 $\pm$ 0.9	24.8 $\pm$ 2.9	2963 $\pm$ 626	1166 $\pm$ 264	2515 $\pm$ 67	768 $\pm$ 63

## Experiments – Result

- Our methods achieve the highest success rates across all test tasks on the ML-10 benchmark.
- However all methods are challenged by the tasks “Shelf-place” and “Lever-pull”.
  - They include unseen objects not included in the raw observation.
  - These tasks cannot be decomposed into previously seen subtasks but rather require new elementary subtasks.
- Link to demo videos: <https://sites.google.com/view/sdvt-neurips>.

Table 3: Meta-World V2 ML-10 success rates on test tasks.

Index. Task	SDVT	SDVT-LW	SD	SD-LW	RL <sup>2</sup>	MAML	PEARL	VariBAD	LDM
11. Drawer-open	<b>65.0±19.9</b>	30.5±12.9	48.8±23.4	45.0±24.0	2.2±1.9	15.8±19.3	1.5±1.1	12.8±12.8	21.8±12.0
12. Door-close	7.5±9.0	<b>81.2±19.0</b>	33.8±25.4	18.8±24.1	8.2±7.3	3.2±6.0	1.2±1.5	27.0±21.8	30.2±28.2
13. Shelf-place	0.0±0.0	<b>1.0±1.2</b>	0.0±0.0	0.0±0.0	0.2±0.2	0.0±0.0	0.0±0.0	0.0±0.0	0.0±0.0
14. Sweep-into	<b>90.0±8.5</b>	51.2±18.9	71.2±14.9	55.0±21.6	64.5±8.3	0.0±0.0	0.8±1.0	30.5±22.0	46.5±21.8
15. Lever-pull	1.2±2.3	3.2±2.7	0.0±0.0	<b>12.5±10.2</b>	0.5±0.8	0.5±0.9	0.5±0.9	0.2±0.5	0.5±0.9
Test mean	32.8±3.9	<b>33.4±5.0</b>	30.8±7.7	26.2±8.7	15.1±2.7	3.9±3.7	0.8±0.5	14.1±6.1	19.8±6.0

## Limitation

---

- Test tasks involving entirely novel subtasks
  - SDVT can prepare for unseen compositions of **seen subtasks**.
  - However, SDVT does not prepare for unseen compositions of **unseen subtasks**.
- Limited to scenarios with **varying rewards**
  - SDVT does not consider setups where **transition dynamics, action, and state space** may vary.
- No temporal information of subtasks
  - Currently the subtask composition is a belief about how the current meta-episode's task is decomposed into.
  - It does not consider **temporal information** of subtasks, such as the ordering or start and termination of subtasks.



# Conclusion

---

1. Latent Dynamics Mixture (LDM): Parametric task variability
  - A robust meta-RL algorithm that effectively prepares for potential **out-of-distribution** test tasks.
  - LDM pretrains the policy on **imaginary tasks** generated from mixtures of latent beliefs on training tasks.
  - LDM effectively enhances generalization to unseen tasks without additional test-time training.
2. Subtask Decomposition and Virtual Training (SDVT): Non-parametric task variability
  - SDVT generalizes LDM for scenarios with **non-parametric task variability**.
  - SDVT employs a Gaussian mixture VAE to meta-learn the set of **elementary subtasks** and the **composition** of each task.
  - SDVT extends the idea of virtual training from LDM to generate tasks with **unseen compositions** of subtasks.
3. Both methods demonstrate the efficacy of virtual training with generated **imaginary tasks**.
  - We may extend the ideas to more general scenarios with varying transition dynamics, state space, or action space.
  - We may combine it with orthogonal strategies such as offline RL or test-time adaptation techniques.

---

Thank you!

# Appendix

## Bayes-Adaptive Belief Update

---

- Posterior belief update

$$b_{t+1} \left( R^{(i)}, T^{(i)} \right) = \frac{\mathbb{P} \left( s_{t+1}, r_{t+1} \mid s_t, a_t, R^{(i)}, T^{(i)} \right)}{\sum_{k=1}^m \mathbb{P} \left( s_{t+1}, r_{t+1} \mid s_t, a_t, R^{(k)}, T^{(k)} \right) b_t \left( R^{(k)}, T^{(k)} \right)} b_t \left( R^{(i)}, T^{(i)} \right). \quad (26)$$

$$\mathbb{P} \left( s_{t+1}, r_{t+1} \mid s_t, a_t, R^{(i)}, T^{(i)} \right) b_t \left( R^{(i)}, T^{(i)} \right) = \mathbb{P} \left( s_{t+1}, r_{t+1} \mid \tau_{:t}, R^{(i)}, T^{(i)} \right) \mathbb{P} \left( R^{(i)}, T^{(i)} \mid \tau_{:t} \right) \quad (27)$$

$$= \frac{\mathbb{P} \left( s_{t+1}, r_{t+1}, \tau_{:t}, R^{(i)}, T^{(i)} \right) \mathbb{P} \left( R^{(i)}, T^{(i)}, \tau_{:t} \right)}{\mathbb{P} \left( R^{(i)}, T^{(i)}, \tau_{:t} \right) \mathbb{P} \left( \tau_{:t} \right)} \quad (28)$$

$$= \frac{\mathbb{P} \left( \tau_{:t+1}, R^{(i)}, T^{(i)} \right)}{\mathbb{P} \left( \tau_{:t} \right)} \quad (29)$$

$$= \mathbb{P} \left( R^{(i)}, T^{(i)} \mid \tau_{:t+1} \right) \frac{\mathbb{P} \left( \tau_{:t+1} \right)}{\mathbb{P} \left( \tau_{:t} \right)} \quad (30)$$

$$= b_{t+1} \left( R^{(i)}, T^{(i)} \right) \mathbb{P} \left( s_{t+1}, r_{t+1} \mid \tau_{:t} \right). \quad (31)$$

## SDVT ELBO Derivation

---

- GMVAE ELBO Derivation

$$\log p_\theta(\tau_{:H^+}) = \log \mathbb{E}_{q_\phi(y_t, z_t | h_t)} \left[ \frac{p_\theta(\tau_{:H^+}, y_t, z_t)}{q_\phi(y_t, z_t | h_t)} \right] \quad (32)$$

$$\geq \mathbb{E}_{q_\phi(y_t, z_t | h_t)} \left[ \log \frac{p_\theta(\tau_{:H^+}, y_t, z_t)}{q_\phi(y_t, z_t | h_t)} \right] \quad (33)$$

$$= \mathbb{E}_{q_\phi(y_t, z_t | h_t)} \left[ \log \frac{p_\theta(\tau_{:H^+} | y_t, z_t) p_\theta(z_t | y_t) p(y_t)}{q_\phi(y_t | h_t) q_\phi(z_t | h_t, y_t)} \right] \quad (34)$$

$$= \mathbb{E}_{q_\phi(y_t, z_t | h_t)} \left[ \log p_\theta(\tau_{:H^+} | z_t) + \log \frac{p_\theta(z_t | y_t)}{q_\phi(z_t | h_t, y_t)} + \log \frac{p(y_t)}{q_\phi(y_t | h_t)} \right]. \quad (35)$$

Equation (35) is equivalent to the ELBO objective in Equation (17) without weighting coefficients. We assume that the reconstruction  $\tau_{:H^+}$  is conditionally independent of the subtask composition  $y_t$  given  $z_t$ .

## Dropout Ablation – LDM

- Dropout rate of LDM
  - Upward trajectory in test performance in correlation with an increase in the dropout rate.
  - However, this trend reverses when the rate approaches its upper limit, specifically at  $p_{\text{drop}} = 0.9$ .




Figure 33: Varying dropout rate. Results of LDM with different dropout rates on Gridworld task.

## Dropout Ablation – Baselines

- We assess the performance of both RL<sup>2</sup> dropout and VariBAD dropout,
  - Dropout destabilizes policy training with multi-step policy gradient loss
  - Dropout on the decoder input improves generalization due to the relatively simpler single-step regression loss with the buffer updates.




Figure 34: Dropout ablation. Results of LDM, RL<sup>2</sup>, and VariBAD trained with and without dropout on Gridworld task.

## SDVT Ablation

- The default setup of our methods are
  - SDVT:  $K = 10$ ,  $\alpha_c = 1.0$ ,  $\alpha_0 = 1.0$ ,
  - SDVT-LW:  $K = 5$ ,  $\alpha_c = 0.5$ ,  $\alpha_0 = 0.0$ .
- When  $\alpha_c$  is too small, task classification collapses into a few one-hot subtasks.
- When  $\alpha_c$  is too large, all tasks return a uniform probability distribution over subtasks.




(a) SDVT ( $K = 10$ ,  $\alpha_c = 1.0$ ,  $\alpha_o = 1.0$ ).(b) SDVT-LW ( $K = 5$ ,  $\alpha_c = 0.1$ ,  $\alpha_o = 0.0$ ).(c) SDVT-LW ( $K = 5$ ,  $\alpha_c = 2.0$ ,  $\alpha_o = 0.0$ ).

Figure 35: Learned subtask compositions on ML-10 with different hyperparameters.



## SDVT Ablation

---

- We report the mean success rates by changing hyperparameters from the default SDVT-LW ( $K = 5, \alpha_c = 0.5, \alpha_o = 0.0$ ) on ML-10.
- We report the difference caused by the changes in parenthesis.

Table 4: Ablation results of SDVT-LW. Performance measured as test success rate (%) on ML-10.

SDVT-LW	without		$K$			$\alpha_c$		
	Dropout	Dispersion	3	7	10	0.1	1.0	2.0
ML-10 Train	65.6	73.0	60.1	69.5	70.8	59.8	70.3	66.3
	(+3.5)	(+10.9)	(-2.0)	(+7.4)	(+8.7)	(-2.3)	(+8.2)	(+4.2)
ML-10 Test	16.5	21.5	25.0	22.0	21.1	20.5	16.1	17.9
	(-16.9)	(-11.9)	(-8.4)	(-11.4)	(-12.3)	(-12.9)	(-17.3)	(-15.5)

## SDVT Learned Compositions – Different Seeds

- Decomposition processes differ among random seeds due to varying initialization and sample tasks during meta-training.
- One subtask can be interpreted as a combination of multiple subtasks on other seeds.
- We rarely have a collapse to a single Gaussian.

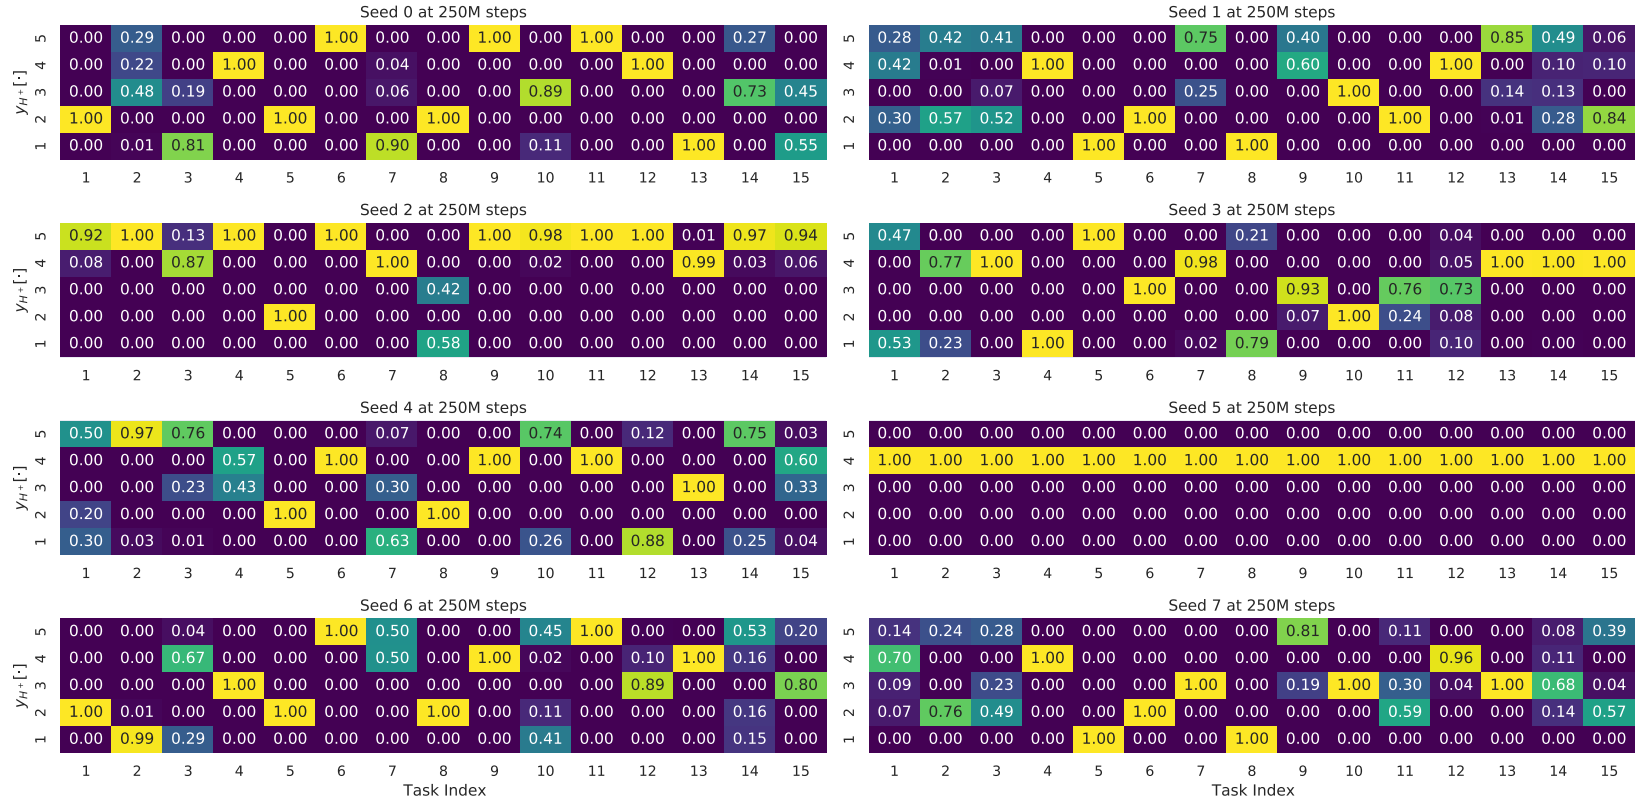


Figure 36: Subtask compositions of all seeds. We visualize the subtask compositions of SDVT-LW on ML-10 after 250M training steps.

## Learned Compositions – Over Time

- Learned compositions over the course of training (0M  $\sim$  250M steps).
- The composition starts with a uniform distribution for all tasks.
- As training progresses, the agent learns to merge similar tasks.




Figure 37: Subtask compositions learned over training. SDVT-LW ( $K = 5, \alpha_c = 0.5, \alpha_o = 0.0$ ) on ML-10.

## Number of Parameters

- We report the number of parameters used by our methods and baselines.
- We demonstrate that our gain is not mainly from the increased capacity.
  - VariBAD (hidden  $\times 2$ ) with hidden dimensions of 512 possesses more parameters than SDVT.
  - We add MLP layers with hidden sizes of [1600, 256] into the encoder and [128] into the policy. The decoder’s hidden size is increased from [64, 64, 32] to [128, 256, 160] to match the capacities of VariBAD (matched) and SDVT-LW.

Table 5: Number of parameters and success rates on ML-10.

Methods	Number of Parameters				ML-10 Success Rate (%)	
	Encoder	Decoder	Policy	Sum	Train	Test
SDVT-LW	1,047,455	174,821	235,401	1,457,677	62.1	<b>33.4</b>
SD-LW	1,047,455	25,637	235,401	1,308,493	<b>75.5</b>	26.2
LDM	502,580	25,577	202,249	730,406	56.7	19.8
<b>LDM (matched)</b>	2,144,692	175,593	267,401	2,587,686	64.2	22.0
VariBAD	251,290	25,577	202,249	479,116	58.2	14.1
<b>VariBAD (hidden <math>\times 2</math>)</b>	894,362	25,577	663,305	1,583,244	62.4	12.0
<b>VariBAD (matched)</b>	1,072,346	175,593	267,401	1,515,340	67.6	17.2

## Computational Complexity

- We report the total wall-clock time required to generate the results for the Half-cheetah-velocity.
  - LDM requires more time than  $RL^2$ , due to more complex architecture that involves the simultaneous training of both the policy network and an independent latent dynamics network.

Table 6: Computation complexity on Half-cheetah-velocity.

	LDM	Mixreg	$RL^2$	VariBAD	ProMP	E-MAML	PEARL
Half-cheetah-velocity Runtime (hours)	31	28	25	10	2	2	25

- We report the total wall-clock time required to generate the results for the ML-10,
  - Despite incorporating the GMVAE and virtual training, our method’s computational demand does not substantially surpass that of VariBAD.

Table 7: Computational complexity on ML-10.

	SDVT	SDVT-LW	SD	SD-LW	$RL^2$	MAML	PEARL	VariBAD	LDM
Wall-clock time (hours)	142	140	138	135	192	17	258	126	131